

Local LLM Document Summarizer Test

Overview:

This is a test document for the Local LLM-powered PDF and HTML summarization system. The system has the following features:

- Works on low-spec GPU or CPU-only environments
- Utilizes open-source LLM models
- Privacy-focused (no external data transmission)
- Supports PDF, HTML, and URL inputs

Technical Features:

The main technical components of the system are:

Document Processing:

PDF files use PyPDF2 or pdfminer.six for text extraction, while HTML files use BeautifulSoup4 for text extraction.

Natural Language Processing:

Extracted text is summarized by locally executable LLM models such as llama.cpp or GPT4All.

Performance Goals:

The system sets the following performance targets:

- Small PDF (1-10 pages): within 3 minutes
- Medium PDF (11-50 pages): within 10 minutes
- Large PDF (51-100 pages): within 30 minutes
- Memory usage: under 8GB during processing

This document was created for testing purposes.